

Reinforcement Learning & Text Games

Hadi Vafaei

0) Problem and motivation

- Reinforcement Learning has been successful in playing Atari games; however, not as much progress has been made for text-based games due to additional language challenges
- Our goal is to build RL agents that can play text-based games. For this study we chose to use TextWorld. Example screenshot from the game →
- At each stage of the game:
 - the agent is provided with a text description of the world state (**obs**)
 - agent plays by generating text **commands** based on these observations
 - sometimes useful commands lead to **rewards**
- These reward signals can be used to train the agent
- The goal is to accumulate high rewards in little time

Goal

You are hungry! Let's cook a delicious meal. Check the cookbook in the kitchen for the recipe. Once done, enjoy your meal!

Obs

-- Kitchen --
Ah, the kitchen. This is some kind of kitchen, really great ordinary vibes in this place, a wonderful ordinary atmosphere.

Cmds

You make out a fridge. Were you looking for an oven? Because look over there, it's an oven. You can't wait to tell the folks at home about this! You see a table. The table is massive. But there isn't a thing on it. You see a counter. You see a cookbook and a knife on the counter. You make out a stove. But the thing is empty, unfortunately. Hm. Oh well

> open fridge
> open fridge
You open the fridge, revealing a yellow bell pepper, a raw yellow potato and a black pepper.

> get yellow bell pepper
> get yellow bell pepper
You take the yellow bell pepper from the fridge.

Reward

Your score has just gone up by one point.

1) Challenges

- Challenges on the language side include:
 - Natural language understanding
 - Text generation
 - Dealing with combinatorial action space

We provide the agent with the list of admissible commands to circumvent some of these difficulties for now...

- RL challenges include:
 - Sparsity of reward
 - Credit assignment
 - Exploration/exploitation

2) Technical Approach

- GRUs + Advantage Actor Critic (A2C): a policy gradient based learning algorithm

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) v_t] \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^{\pi}(s, a)] \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi}(s, a)] \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \delta]\end{aligned}$$

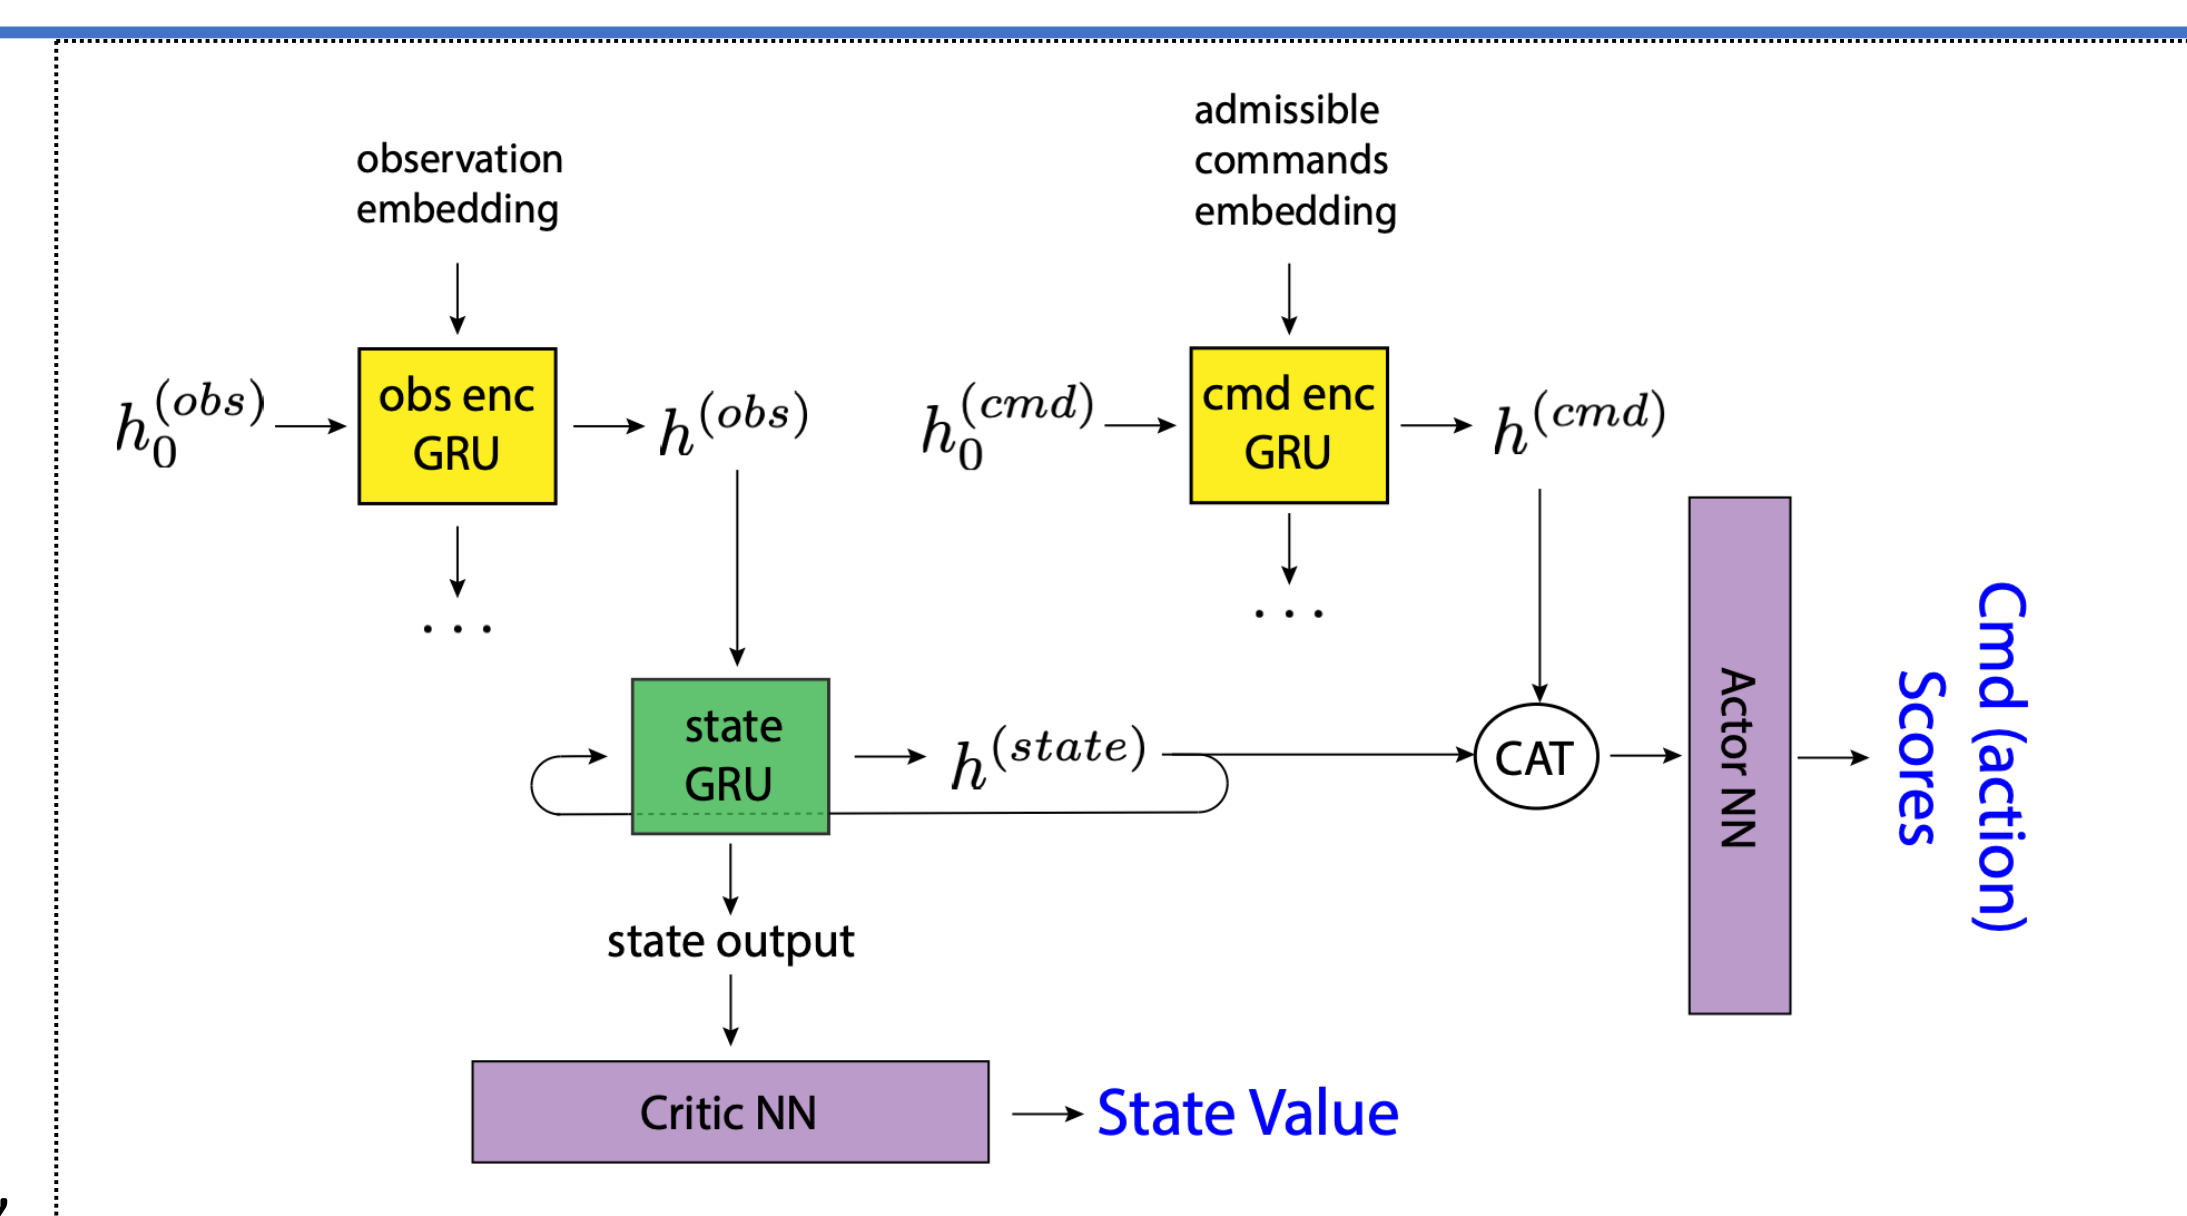
REINFORCE
Q Actor-Critic
Advantage Actor-Critic
TD Actor-Critic

→ This is what we used

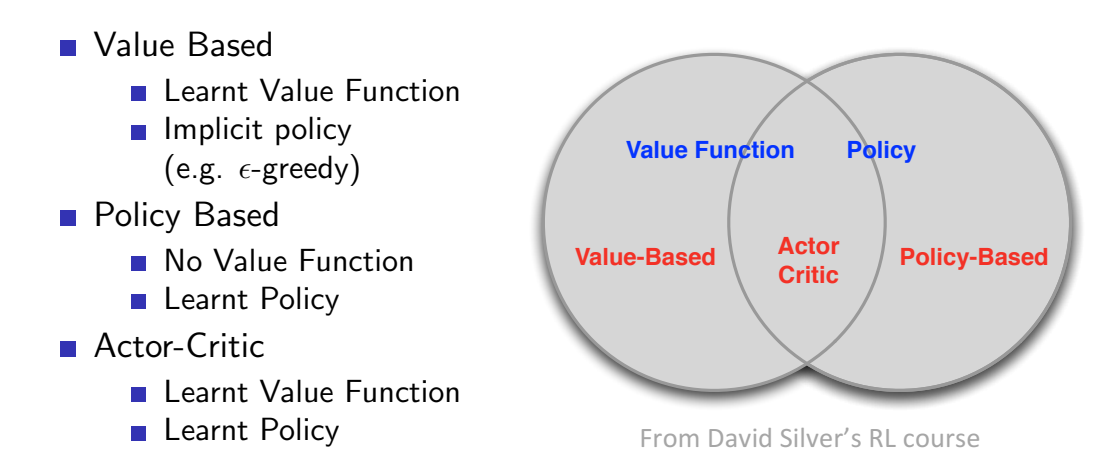
- Inspired by Soft Actor Critic algorithm (SAC), we experimented with adding an entropy term to our loss to see if it helps (it doesn't)

- Advantage function is defined as:

$$A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$$




- Discrete actions: Softmax Policy



3) Results


Experiment # 1 (simple games):

- Trained on 100 games with dense rewards and a detailed description of goal
- Validated on 20 games from similar distribution
- Observations:
 - Our agent consistently outperforms the random baseline by a large margin but is inferior to humans
 - Including the entropy term hurts model performance. Perhaps this is because SAC is better suited for tasks with continuous action space



Experiment # 2 (hard games):

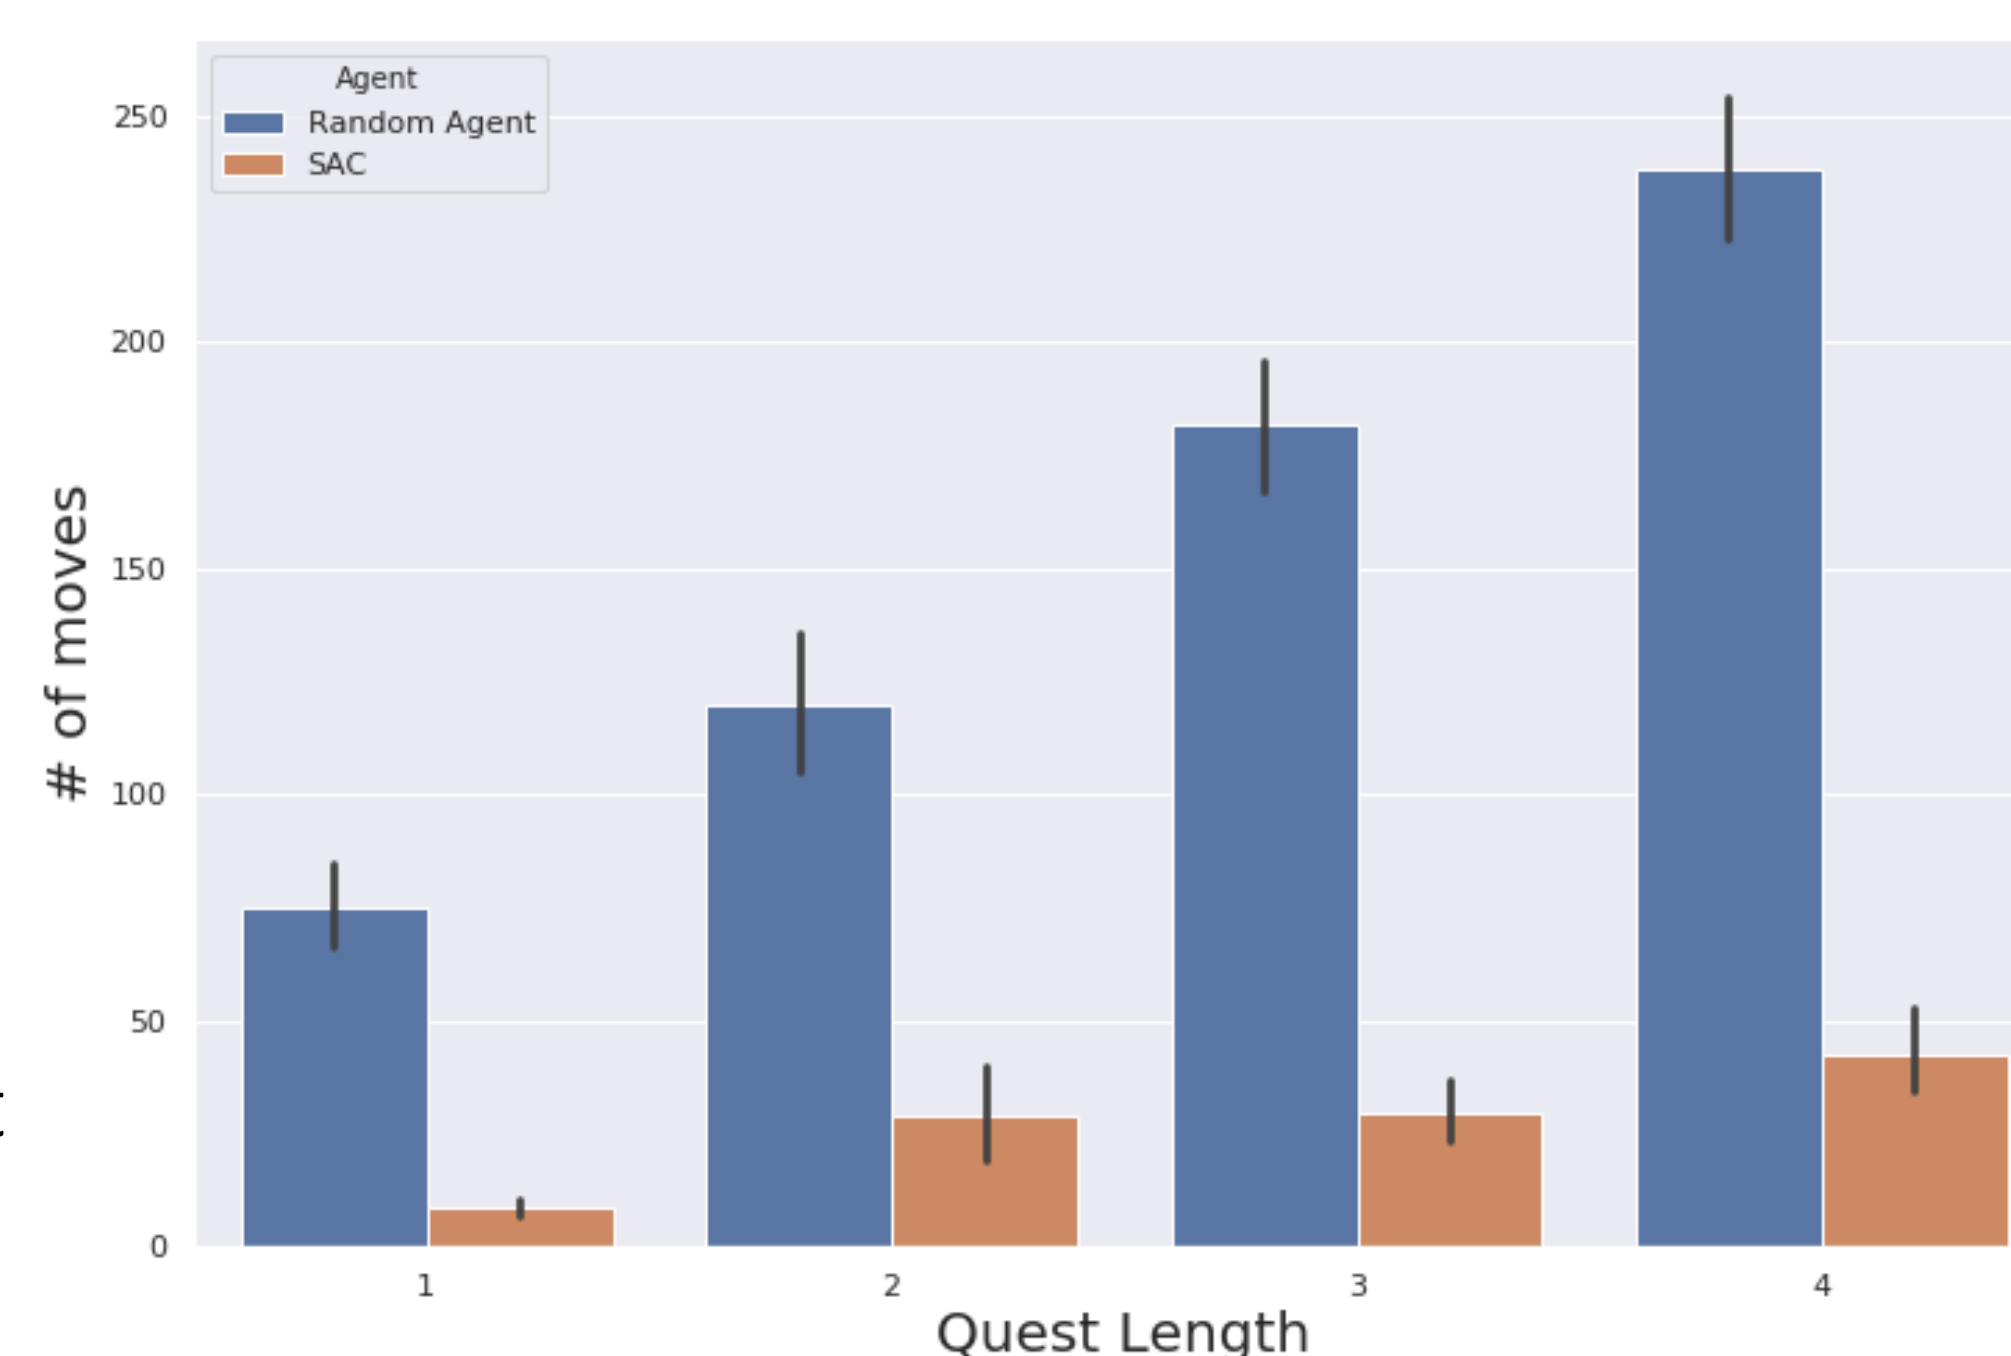
- We trained the same model on 4440 games from the competition
- Validated on 222 games
- Due to difficulty of the tasks, our agent barely outperforms the random baseline. This hints at the need for better architecture/memory



Experiment # 3 (quest length):

- Trained/validated on games with quest length of {1, 2, 3, 4}
- For example, quest length of 1 means the agent instantly wins the game at the first stage if it makes the right choice

- Observations:
 - Performance much better than the random agent
 - These games are simple, so humans will win a game of quest length N in roughly ~N steps
 - Therefore, our agent is again inferior to humans



Here lambda is the entropy term hyperparameter in the loss:

$$Loss = policy\ loss + value\ loss + \lambda_H * H$$

4) Future Work

- Meta-Learning: better than joint training on batch of games?
- Some sort of supervision: Imitation Learning?
- Better model architectures, structured memory

Key References:

- TextWorld: A Learning Environment for Text-based Games [Côté et. al. 2018]
- Language Understanding for Text-based Games using Deep Reinforcement Learning [Narasimhan et. al. 2015]
- Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor [Haarnoja et. al. 2018]